

# COMP9120

Week 12: Review

Semester 2, 2022

Professor Athman Bouguettaya  
School of Computer Science



THE UNIVERSITY OF  
SYDNEY



# Acknowledgement of Country

*I would like to acknowledge the Traditional Owners of Australia and recognise their continuing connection to land, water and culture. I am currently on the land of the Darug people and pay my respects to their Elders, past, present and emerging.*

*I further acknowledge the Traditional Owners of the country on which you are on and pay respects to their Elders, past, present and future.*

---



## COMMONWEALTH OF AUSTRALIA

### Copyright Regulations 1969

#### **WARNING**


This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**).

The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

**Do not remove this notice.**




Average: 48%





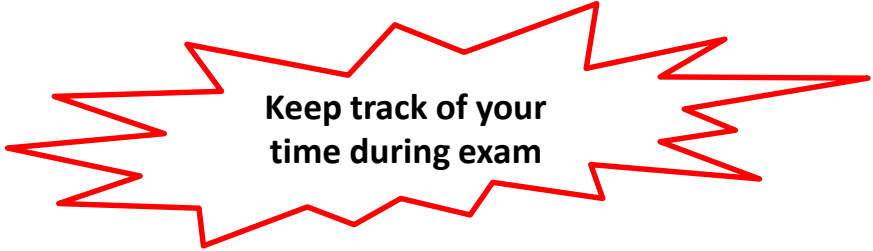
# SQL Challenge Summary

Average: 92%





- Date: Wednesday 16th November
- Time: 1:00 PM Sydney time
- Duration: 175 Minutes
  - Reading time: 10 Minutes
  - Writing time: 150 Minutes
  - Upload time: 15 minutes
- Everyone starts the exam at the same time
  - No late submission
- Exam adjustment is done by the exam office
  - Notification no later than 3 days before the exam



**Keep track of your  
time during exam**

## Question Type

- No MCQs
- Essay-type questions (10 Questions → 50 marks)
  - ERD
  - Relational Model
  - Relational Algebra
  - SQL
  - Integrity Constraints
  - Transactions
  - Normalizations
  - Storage
  - Indexing
  - Query Processing



Practice exam is  
available in Canvas

- › Exam covers 50% of the final mark
- › You must obtain at least **40% in the exam**, as well as an **overall mark of at least 50%**, to pass the unit



	Week	Topic
Foundations	Week 1	Introduction
	Week 2	Conceptual Database Design
	Week 3	Relational Data Model / Logical Database Design
	Week 4	Relational Algebra and SQL
	Week 5	Complex SQL
	Week 6	Database Integrity
Applications	<del>Week 7</del>	<del>Database Application Development and Security</del>
	Week 8	Transaction Management
	Week 9	Schema Refinement and Normalisation
Internals	Week 10	Storage and Indexing
	Week 11	Query Evaluation and Optimisation
	Week 12	Revision
	Week 13	Final exam Structure





```
CREATE TABLE Film(  
  filmID int PRIMARY KEY,  
  title varchar(30),  
  releaseYear int);
```

```
CREATE TABLE Actor(  
  actorID int PRIMARY KEY,  
  actorName varchar(30),  
  nationality varchar(20));
```

```
CREATE TABLE FilmActor(  
  filmID int,  
  actorID int DEFAULT 0,  
  PRIMARY KEY(filmID, actorID),  
  FOREIGN KEY (filmID) REFERENCES Film ON DELETE CASCADE,  
  FOREIGN KEY (actorID) REFERENCES Actor ON UPDATE SET DEFAULT);
```

Film		
filmID	title	releaseYear
101	No time to die	2021
102	Titanic	1998
205	Baby Day out	2000

Actor		
actorID	actorName	nationality
0	No Actor	null
981	Daniel	UK
965	Kate	USA
901	Lily	Australia

FilmActor	
filmID	actorID
101	981
102	965
205	901

What will be the records of Film and FilmActor table after executing the following query?

**DELETE FROM Film WHERE filmID = 101;**

Film		
filmID	title	releaseYear
102	Titanic	1998
205	Baby Day out	2000

FilmActor	
filmID	actorID
102	965
205	901



```
CREATE TABLE Film(  
  filmID int PRIMARY KEY,  
  title varchar(30),  
  releaseYear int);
```

```
CREATE TABLE Actor(  
  actorID int PRIMARY KEY,  
  actorName varchar(30),  
  nationality varchar(20));
```

```
CREATE TABLE FilmActor(  
  filmID int,  
  actorID int DEFAULT 0,  
  PRIMARY KEY(filmID, actorID),  
  FOREIGN KEY (filmID) REFERENCES Film ON DELETE CASCADE,  
  FOREIGN KEY (actorID) REFERENCES Actor ON UPDATE SET DEFAULT);
```

Film		
filmID	title	releaseYear
101	No time to die	2021
102	Titanic	1998
205	Baby Day out	2000

Actor		
actorID	actorName	nationality
0	No Actor	null
981	Daniel	UK
965	Kate	USA
901	Lily	Australia

FilmActor	
filmID	actorID
101	981
102	965
205	901

What will be the records of Actor and FilmActor table after executing the following query?

**UPDATE** Actor **SET** actorID = 999 **WHERE** actorID = 981;

Actor		
actorID	actorName	nationality
0	No Actor	null
999	Daniel	UK
965	Kate	USA
901	Lily	Australia

FilmActor	
filmID	actorID
101	0
102	965
205	901



```
CREATE TABLE Film(  
    filmID int PRIMARY KEY,  
    title varchar(30),  
    releaseYear int);
```

```
CREATE TABLE Actor(  
    actorID int PRIMARY KEY,  
    actorName varchar(30),  
    nationality varchar(20));
```

```
CREATE TABLE FilmActor(  
    filmID int,  
    actorID int DEFAULT 0,  
    PRIMARY KEY(filmID, actorID),  
    FOREIGN KEY (filmID) REFERENCES Film ON DELETE CASCADE,  
    FOREIGN KEY (actorID) REFERENCES Actor ON UPDATE SET DEFAULT);
```

› Create an assertion for the following:

- An actor cannot take part in more than 3 movies in a year

```
CREATE ASSERTION ActingConstraint CHECK (  
    NOT EXISTS (  
        SELECT 1  
        FROM Film NATURAL JOIN FilmActor  
        GROUP BY actorID, releaseYear  
        HAVING COUNT(filmID) > 3  
    )  
);
```

Film		
filmID	title	releaseYear
101	No time to die	2021
102	Titanic	1998
205	Baby Day out	2000

Actor		
actorID	actorName	nationality
0	No Actor	null
981	Daniel	UK
965	Kate	USA
901	Lily	Australia

FilmActor	
filmID	actorID
101	981
102	965
205	901

- › Consider the following relational schema:

*Emp*(*eid*, *ename*, *age*, *salary*)

*Works*(*eid*, *did*, *since*)

*Dept*(*did*, *budget*, *managerid*)

Define an assertion that will ensure that all managers have *age* > 30.

- › **CREATE ASSERTION** managerAge

**CHECK** ( **NOT EXISTS** (**SELECT** age

**FROM** Emp, Dept

**WHERE** eid = managerid **AND** age <= 30))

Student(snum, sname, major, level, age)

Class(name, meets\_at, room, fid)

Enrolled(snum, cname)

Faculty(fid, fname, deptid)

```
CREATE TABLE Enrolled (snum INTEGER, cname CHAR(20),
    PRIMARY KEY (snum, cname),
    FOREIGN KEY (snum) REFERENCES Student,
    FOREIGN KEY (cname) REFERENCES Class,
    CHECK (( SELECT COUNT (snum)
              FROM Enrolled
              GROUP BY cname) <= 30))
```

Express each of the following integrity constraints in SQL


- › Every class has a maximum enrolment of 30 students.
- › Every faculty member must teach at least two courses.
- › No department can have more than 10 faculty members.

```
CREATE ASSERTION TeachConstraint
CHECK ( ( SELECT COUNT (*)
          FROM Faculty F NATURAL JOIN Class C
          GROUP BY C.fid
          HAVING COUNT (*) < 2) = 0)
```

```
CREATE TABLE Faculty (fid INTEGER, fname CHAR(20), deptid INTEGER,
    PRIMARY KEY (fnum),
    CHECK ( ( SELECT COUNT (*)
              FROM Faculty F
              GROUP BY F.deptid
              HAVING COUNT (*) > 10) = 0))
```

Determine whether the following schedule is conflict serializable; justify your answer. If it is conflict serializable, please give a conflict equivalent serial schedule.

›  $R1(x), W2(y), R1(z), R3(z), W2(x), R1(y)$



There are two conflicts in this case.

First conflict is between  $R1(x)$  and  $W2(x)$ .

Because of this conflict we need to put T1 before T2 i.e  $T1 \rightarrow T2$ .

Second conflict is between  $W2(y)$  and  $R1(y)$ .


Because of this conflict we need to put T2 before T1 i.e  $T2 \rightarrow T1$

Since we have a cycle between T1 and T2, that's why this schedule is NOT conflict serializable.

---

Determine whether the following schedule is conflict serializable; justify your answer. If it is conflict serializable, please give a conflict equivalent serial schedule.

›  $R1(x), W2(y), R1(z), R3(x), W2(x), R2(y)$



There are two conflicts in this case.

First conflict is between  $R1(x)$  and  $W2(x)$ .

Because of this conflict we need to put T1 before T2 i.e  $T1 \rightarrow T2$ .

Second conflict is between  $R3(x)$  and  $W2(x)$ .

Because of this conflict we need to put T3 before T2 i.e  $T3 \rightarrow T2$

Since there is no cycle, so this schedule is conflict serializable.


According to the condition, T1 and T3 must be performed before T2. So, either of T1, T3, T2 OR T3, T1, T2 is a conflict equivalent serial schedule in this case.

---

› Given the following schedule S:

$W1(x), R1(x), W3(z), R2(x), R2(y), R4(z), R3(x), W3(y)$

Check whether this schedule is conflict serializable. If the schedule is conflict equivalent to a serial schedule, choose the equivalent serial order.



There are 4 conflicts in this case.

First conflict is between  $W1(x)$  and  $R2(x)$ . Because of this conflict we need to put T1 before T2 i.e  $T1 \rightarrow T2$ .

Second conflict is between  $W1(x)$  and  $R3(x)$ . Because of this conflict we need to put T1 before T3 i.e  $T1 \rightarrow T3$ .

Third conflict is between  $W3(z)$  and  $R4(z)$ . Because of this conflict we need to put T3 before T4 i.e  $T3 \rightarrow T4$ .

Fourth conflict is between  $R2(y)$  and  $W3(y)$ . Because of this conflict we need to put T2 before T3 i.e  $T2 \rightarrow T3$ .

Since there is no cycle, so this schedule is conflict serializable. The equivalent serial order is T1, T2, T3, T4.



Consider a database with a table Results(StudentId, UnitCode, Grade) that contains.

(Row)	StudentId	UnitCode	Grade
A	3245	INFO2120	57
B	3245	COMP2129	82
C	4290	INFO2120	68
D	4290	MATH2002	56

The steps in Figure 1 can be treated for the purposes of transaction management as an execution schedule consisting of read (r) and write (w) operations by either transaction 1 or 2 upon different objects - in this case we shall assume that the granularity of these objects is row-level, with each row denoted by the letter given in the table above.

Which of the following execution schedule matches the steps in Figure 1?

1. r1(B),r2(A),r2(B),r2(C),r2(D),r1(A),r1(B),r1(A),r1(C),r2(A),r2(C)
2. r1(B),r2(A),r2(B),w2(C),w2(D),r1(A),r1(B),r1(A),r1(C),r2(A),r2(C)
- ✓ 3. r1(B),r2(A),r2(B),r2(C),w2(C),r2(D),w2(D),r1(A),r1(B),r1(A),r1(C),r2(A),r2(C)

T1:	Begin Transaction
T1:	SELECT Grade FROM Results WHERE StudentId = '3245' AND UnitCode = 'COMP2129'
T2:	Begin Transaction
T2:	SELECT UnitCode, Grade FROM Results WHERE StudentId = '3245'
T2:	UPDATE Results SET Grade = Grade + 5 WHERE StudentId = '4290'
T1:	SELECT UnitCode, Grade FROM Results WHERE StudentId = '3245'
T1:	SELECT Sum(Grade) FROM Results WHERE UnitCode = 'INFO2120'
T1:	COMMIT
T2:	SELECT Sum(Grade) FROM Results WHERE UnitCode = 'INFO2120'
T2:	COMMIT

Figure 1



Which FD might be valid?

A	B	C	D
1	1	2	a
1	2	1	b
2	2	3	c
1	2	3	d
3	2	1	e
1	2	1	f

1.  ~~$AB \twoheadrightarrow C$~~

✓ 2.  $AC \rightarrow B$

3.  ~~$BC \rightarrow A$~~

4.  ~~$A \rightarrow BC$~~

5.  ~~$B \twoheadrightarrow AC$~~

- › Consider the following instance of a relation.

A	B	C	D
1	x	a	10
2	y	b	20
3	z	b	20
1	z	c	30
2	y	b	20

- › Which of the following functional dependencies might be valid for this relation?

✓ 1.  $AB \rightarrow C$

✓ 2.  $C \rightarrow D$

3.  $A \rightarrow B$

✓ 4.  $ABC \rightarrow D$

5.  $B \rightarrow C$

✓ 6.  $D \rightarrow C$

Consider attributes of the relation, describing the cost of damage to people who have been involved in car accidents:

Damage(regno, license, address, accidentID, amount)

Match each of the business rules below with the corresponding functional dependency.

- |    |  |          |
|----|--|----------|
| A. | A car (regno) can have at most one driver (license)      | <b>4</b> |
| B. | Drivers must be registered at a single address           | <b>5</b> |
| C. | A driver can be in at most one car in any given accident | <b>6</b> |

Functional Dependencies

1. regno, accidentID  $\rightarrow$  license
2. license  $\rightarrow$  regno
3. address  $\rightarrow$  license
4. regno  $\rightarrow$  license
5. license  $\rightarrow$  address
6. license, accidentID  $\rightarrow$  regno



$R(A, B, C, G, H, I)$

$F = \{$

$A \rightarrow B$

$A \rightarrow C$

$CG \rightarrow H$

$CG \rightarrow I$

$B \rightarrow H$

$\}$

Is  $(AG)$  a candidate key?

› Compute the attribute closure  $(AG)^+$

1.  $result = AG$
2.  $result = ABG (A \rightarrow B)$
3.  $result = ABCG (A \rightarrow C)$
4.  $result = ABCGH (CG \rightarrow H)$
5.  $result = ABCGHI (CG \rightarrow I)$

› Is  $(AG)$  a candidate key?

1. Is  $(AG)$  a super key? YES! The closure  $(AG)^+ = R$  (i.e., **all** attributes in  $R$ )
2. Is  $(AG)$  minimal? (Is any subset of  $(AG)$  a superkey?)
  - $(A)^+ = ABCH \neq R$
  - $(G)^+ = G \neq R$

*Therefore  $AG$  is a candidate key*

- › Consider a relation  $R(A, B, C, D, E, F)$  with functional dependencies:

$A \rightarrow BDF$

$BC \rightarrow DEF$

$D \rightarrow CE$ .

- › Is  $AB$  a candidate key?

- › Compute the closure  $(AB)^+$

1.  $result = AB$
2.  $result = ABDF (A \rightarrow BDF)$   
 $= ABCDEF (D \rightarrow CE)$

- ›  $AB$  is a superkey because the closure is  $R$

- › Compute the closure  $A^+$

1.  $result = A$
2.  $result = ABDF (A \rightarrow BDF)$   
 $= ABCDEF (D \rightarrow CE)$

- ›  $AB$  is not a candidate key because a subset of  $(AB)$ , i.e.  $A$ , is a key.

Work on a project is tracked with the following relation:

WorksOn (EmpID, ProjID, start, finish, manager)

The following FDs hold over this relation:

EmpID  $\rightarrow$  manager (Each employee can only have one manager)

EmpID, ProjID  $\rightarrow$  start, finish (Each employee can only work on a project once)

ProjID, start  $\rightarrow$  EmpID (Only one employee can start work on the project at a time)

Which of the following are candidate keys for the relation?

- ✓ 1. **ProjID, start**
- ✓ 2. **EmpID, ProjID**
- 3. **start, finish, manager**
- 4. **EmpID, ProjID, start**

Based on the data in the relation, is this a valid MVD?

$UoS \twoheadrightarrow Lecturer$

Note that according to the values in the relation, the relationship between the UoS and Lecturer is **independent** from the relationship between UoS and Textbook. *This means that the above MVD is valid.* This also implies that the Lecturer of a UoS is selected independently by the school.

Assume a new lecturer, Lijun C is added for the UoS COMP9120. What must happen to maintain this independence (i.e., MVD)?

- Add one row for each different textbook with Lecturer *Lijun C* of that UoS.

<u>UoS</u>	<u>Textbook</u>	<u>Lecturer</u>
COMP9120	Silberschatz	Ying Z
COMP9120	Widom	Ying Z
COMP9120	Silberschatz	Mohammad P
COMP9120	Widom	Mohammad P
COMP9120	Silberschatz	Alan F
COMP9120	Widom	Alan F
COMP5110	Silberschatz	Ying Z
COMP5110	Silberschatz	Mohammad P
<b>COMP9120</b>	<b>Widom</b>	<b>Lijun C</b>
<b>COMP9120</b>	<b>Silberschatz</b>	<b>Lijun C</b>



Assume the only key to the following relation is the set (UoS, Textbook, Lecturer).

Is this relation in 4NF?

No: There are at least *two non-trivial multivalued dependencies*

UoS  $\twoheadrightarrow$  Textbook *and* UoS  $\twoheadrightarrow$  Lecturer  
and UoS is *not* a superkey.

Solution: Split the above relation into two relations:

Now both relations are

In 4NF! Why?

MVDs are trivial now!

<u>UoS</u>	<u>Textbook</u>	<u>Lecturer</u>
COMP9120	Silberschatz	Ying Z
COMP9120	Widom	Ying Z
COMP9120	Silberschatz	Mohammad P
COMP9120	Widom	Mohammad P
COMP9120	Silberschatz	Alan F
COMP9120	Widom	Alan F
COMP5110	Silberschatz	Ying Z
COMP5110	Silberschatz	Mohammad P





<b>uosCode</b>	<b><u>lecturerId</u></b>
COMP5138	3456
COMP5338	4567

```
BEGIN;  
UPDATE Course SET lecturerId=4567 WHERE uosCode='COMP5138';  
COMMIT;  
SELECT lecturerId FROM Course WHERE uosCode='COMP5138';
```

1. 1234
- ✓ 2. 3456
3. 4567

- › Give a set of FDs for the relation schema  $R(A,B,C,D)$  with candidate key  $AB$  so that  $R$  is not in 2NF.

2NF: *There cannot be a functional dependency between a subset of a key to non-key attributes.* This is applicable only when the key consists of more than one attribute.

An example of such FDs is:

$$AB \rightarrow CD$$

$$B \rightarrow C$$

$AB$  is obviously a key for this relation since  $AB \rightarrow CD$  implies  $AB \rightarrow ABCD$ . However, the FD  $B \rightarrow C$  violates 2NF since:

$B$  is a proper subset of the key  $AB$

$C$  is not part of a key

- › Give a set of FDs for the relation schema  $R(A,B,C,D)$  with candidate key  $AB$  under which  $R$  is in 2NF but not in 3NF.

### Third Normal Form (3NF)

Formal Definition: a relation  $R$  is in 3NF if for each dependency  $X \rightarrow Y$  in  $F^+$ , *at least one of the following holds*:

$X \rightarrow Y$  is a trivial FD ( $Y \subseteq X$ )

$X$  is a superkey for  $R$

$Y \subset$  (is a proper subset of) a candidate key for  $R$

An example of such FDs is:

$$AB \rightarrow CD$$

$$C \rightarrow D$$

$AB$  is obviously a key for this relation since  $AB \rightarrow CD$  implies  $AB \rightarrow ABCD$ . The FD:  $C \rightarrow D$  violates 3NF but complies with 2NF since:

$C \rightarrow D$  is *not* a trivial FD

$C$  is *not* a superkey

$D$  is *not* part of some key for  $R$

- › Consider a relation  $R$  with attributes  $ABCDE$  and the following FDs:

$$A \rightarrow BC, BC \rightarrow E, \text{ and } E \rightarrow DA.$$

Is  $R$  in BCNF? If not, give a lossless join decomposition of  $R$ .

The schema  $R$  has keys  $A$ ,  $E$  and  $BC$ . Why? Use attribute closure to find keys.

It follows that  $R$  is in BCNF.

- › Let  $S$  be a relation with attributes  $ABCDE$  and the following FDs are given:

$$A \rightarrow CE, D \rightarrow B, \text{ and } E \rightarrow DA.$$

Is  $S$  in BCNF? If not, give a lossless join decomposition of  $R$ .

The schema  $S$  has keys  $A$  and  $E$ . Use attribute closure to find keys.

It follows that  $S$  is *not* in BCNF.

Therefore, decompose  $S$  into  $S_1 = (D, B)$  and  $S_2 = (A, C, D, E)$ . It is lossless join decomposition because the intersection of  $S_1$  and  $S_2$  is a key to  $S_1$ .

Is the following relation in BCNF? If not, give a lossless-join decomposition.

***contracts (contractID, supplierID, projectID, deptID, itemID, quantity, value)***

Functional dependencies:

*contractID*  $\rightarrow$  *supplierID, projectID, deptID, itemID, quantity, value*

*supplierID, deptID*  $\rightarrow$  *itemID*

*projectID*  $\rightarrow$  *supplierID*

**Solution:**

It is not in BCNF. Because in the given 2<sup>nd</sup> and 3<sup>rd</sup> FDs, the LHS is not a key. How do we know?

*Use attribute closure.*

Therefore, looking at ***supplierID, deptID***  $\rightarrow$  ***itemID***, we can divide *contracts* relation into

R1 = (*supplierID, deptID, itemID*) and

R2 = (*supplierID, deptID, contractID, projectID, quantity, value*)

Then looking at ***projectID***  $\rightarrow$  ***supplierID***, we can divide R2 into

R3 = (*projectID, supplierID*) and

R4 = (*projectID, deptID, contractID, quantity, value*)

Therefore, R1, R3 and R4 is the lossless-join decomposition of *contracts* relation. Note the intersection of the above relations is a key to one of the relations.

Assume that there are 1,000,000 records in a relation R and each record is 200 bytes long. Each page is 4K bytes, of which 250 bytes are reserved for header and array of record pointers. Assume pages are 90% full on average.

- › How many records can we store in each page?
- › How many pages are required to store R?

**Solution:**

Empty space in each page is  $(4 * 1024 - 250) = 3846$  bytes

Number of records per page =  $\text{floor}(3846 / 200) = 19$  records

On average, each page contains  $\text{floor}(19 * 90\%) = 17$  records

Number of pages required to store the table =  $\text{ceil}(1,000,000 / 17) = 58,824$  pages

## › True/False?

- Parser takes a query-evaluation plan, executes that plan, and returns the answers. **False**
- SQL systems remove duplicates even if the keyword **DISTINCT** is not specified in a query. **False**
- Pipelined evaluation is cheaper than materialization **True**
- External merge-sort works even if the entire table does not fit in the main memory **True**
- Natural join is a special case of Equi-join **True**



Consider two tables employee(eid, ename, did) and department(did, dname). There are 10000 tuples in the employee table and 1000 tuples in the department table. It requires 500 pages and 100 pages to store the records of employee and department table respectively.

- › Calculate the cost of nested loop join and block-nested loop join of these two tables.

The estimated cost of nested loops join

employee as outer table:  $500 + 10000 * 100 = \mathbf{1,000,500 \text{ disk I/Os}}$

department as outer table:  $100 + 1000 * 500 = \mathbf{500,100 \text{ disk I/Os}}$

The estimated cost of block nested loops join

employee as outer table:  $500 + 500 * 100 = \mathbf{50,500 \text{ disk I/Os}}$

department as outer table:  $100 + 100 * 500 = \mathbf{50,100 \text{ disk I/Os}}$



That's it for today..

See you next week!



THE UNIVERSITY OF  
SYDNEY